Reflection on AI Sycophancy

In my experience, AI sycophancy appears in two forms: **overly rigid adherence to past instructions** and **excessive agreement with the user**.

The first occurs when an AI **turns a temporary request into a lasting rule or a persistent memory**. I might ask it to respond in a certain way for one particular task, but it continues applying that preference in later conversations. Over time, its responses become increasingly shaped by an instruction I never intended to make permanent. This forces me to repeatedly clarify that a request applies only to the next response. Ambiguous prompts contribute to this problem, but the burden should not fall entirely on users.

The second, and more concerning, problem is that AI often **treats a question as a conclusion to endorse**. For example, I might ask, "Is Plan A better than Plan B? Because I think plan A has these advantages: ..." Although my wording suggests a preference, I am still asking for an evaluation. Yet the AI may simply confirm that those advantages are valid and conclude that Plan A is better. It may overlook Plan A's weaknesses and Plan B's strengths. Verifying the reasons I provide is only part of the task; a useful answer requires an independent comparison of both alternatives.

One practical response is to write clear instructions in `AGENTS.md`, explicitly asking the agent to **examine assumptions, reason from first principles, and disagree with or reject my proposals when justified**. However, this cannot eliminate the problem. In my view, AI systems often prioritize completing a user's request over evaluating whether its premises are sound. Users can reduce this tendency by writing clearer prompts and explicitly requesting critical evaluation, although this adds effort. Another approach is to **open a separate conversation with minimal prior context and ask the AI to investigate the question independently**. This can reduce the influence of earlier assumptions and help me reconsider how I framed the problem. I can then bring that analysis back into the main conversation and continue with a better-supported conclusion.